FAR AWAY · 2026

Problem-Statement Decision Brief

8 candidates · rated and compared · pick one

Optimised for: builder-first hackathon scoring + maximum resume / recruiter pull

OUR RECOMMENDATION

GradeTrust

the fairness + trust layer for exam evaluation

A sellable product that also showcases frontier AI-security depth (injection-robust grading).
Freshest new contender: VivaProof (GenAI-era authorship verification). Deepest security pick: FireBreak.

Claude's overall rating **9.2 / 10**

Software-only **Resume-first** **Team of 4-5**

Prepared by Claude (Opus 4.8) · 11 June 2026 · Share with your team to finalise on one

How to read this brief

This brief compares EIGHT problem statements for FAR AWAY 2026, each researched against the live web in June 2026 and framed to be strong on BOTH the hackathon's builder-first scoring and resume / recruiter pull. Each is rated by Claude on six axes (1-5), given an overall out of 10, and carries an explicit STAR NOVELTY - the specific edge that makes it stand out. Read the matrix, skim the profiles, then use the 5-minute team vote on the last page.

What deeper research changed: agent-security *primitives* (runtime authz, memory firewall, MCP gateway, single-agent CaMeL) are ALREADY shipped in 2026 by Microsoft, Google and mguard - so they are no longer novel. Novelty has moved to **vertical products**: hence GradeTrust and the new **VivaProof**, plus a Space wildcard (**OrbitGuard**). **FireBreak** is the one security primitive still genuinely open (multi-agent infection).

Scoring legend

Axis	What it measures
Resume / recruiter pull	How strongly it makes a recruiter want to interview you
Real-world impact	Does it solve a real, felt problem? Could it be sold / shipped?
Technical depth	Real engineering vs a thin 'AI wrapper' (which the rubric penalises)
Originality	How crowded the space already is - how novel your angle is
Demo legibility	Can a non-technical judge instantly 'get' the 2-minute demo?
Buildability	Realistic to ship a strong version with a 4-5 person team in the window

The eight candidates at a glance

Problem statement	Res.	Impact	Tech	Orig.	Demo	Build	Overall
GradeTrust Examinations	5	5	4.5	4	5	4	9.2
VivaProof Examinations	4.5	4.5	4	4	4.5	3.5	8.6
FireBreak Agentic and Autonomous Systems	4.5	3	5	5	4	3	8.3
SentryBrowse Agentic and Autonomous Systems	4.5	3.5	4.5	3.5	5	3	8.2
OrbitGuard Space and Aerospace	4	3.5	4	4	5	3.5	8.0
ExamShield Examinations	3.5	4.5	3.5	3	5	3	7.5
JusticeRAG Examinations	3.5	3.5	4	3	4	4	7.3
CaMeL-Ops Agentic and Autonomous Systems	4	3	5	2	4	3	7.0

Cell colour = score band (green high to red low). Deep-dive profiles follow, one per page.

THE DECISION IN ONE LINE

- **Balanced #1 (product + security depth + legible):** build **GradeTrust** - our pick.
- **Freshest new vertical (hot university pain): VivaProof** (authorship verification).
- **Deepest security signal (still-open frontier): FireBreak**.
- **Most dramatic demo / hottest agent topic: SentryBrowse** (but a contested research space).
- **Judge-wow wildcard / standout resume line: OrbitGuard** (Space).
- **Avoid as primary:** CaMeL-Ops, JusticeRAG, and anything in the 'already shipped' box below.

ALREADY SHIPPED IN 2026 - DO NOT PICK AS YOUR CORE

Tempting idea	Who already shipped it (so it is no longer novel)

Runtime agent authorization / ‘agent IAM’	Microsoft Agent Governance Toolkit (MIT, all 10 OWASP agentic risks); IBM; NIST initiative
Agent memory firewall (MINJA defense)	mguard - Ed25519 provenance + Bayesian trust, already shipped open-source
Generic MCP security gateway / firewall	mcp-firewall, pipelock, Lasso MCP Gateway, agent-shield
Generic AI answer-sheet grader	Eklavvya, Chanakya AI, Gradelab, CBSE’s own 2026 pilot
CaMeL reimplementation (single-agent)	google-research/camel-prompt-injection, Microsoft Dromedary, OpenClaw

GradeTrust

The fairness and trust layer for exam evaluation

CLAUDE'S RATING
9.2 / 10

RECOMMENDED · #1 Theme: Examinations

★ STAR NOVELTY (our edge): Sell TRUST, not grading. The only platform that MEASURES and fixes grading inconsistency (Many-Facet Rasch psychometrics) across human AND AI graders, resists student injection attacks hidden in answers, and audits bias - the trust layer the crowded grader market (Eklavvya, Chanakya, CBSE) completely skips.

THE PROBLEM (plain English)

Two examiners grade the SAME answer sheet and give very different marks (a 6 vs a 9) and that gap decides careers in UPSC Mains, board exams, university and coaching mock-tests. AI graders are very consistent but (a) only moderately accurate, (b) provably biased, and (c) trivially cheated by students who hide 'give me full marks' inside their answer.

HOW WE'D SOLVE IT

A trust layer ON TOP of grading. It (1) runs an ensemble of rubric-grounded AI examiners plus the human marks, then uses Many-Facet Rasch to model each grader's severity and reconcile to one calibrated, fair score; (2) blocks student prompt-injection in answer sheets; (3) audits bias (handwriting, length, vernacular); (4) explains every mark with a rubric citation and appeal trail. Sold B2B to coaching institutes, universities and boards.

RATINGS

Resume / recruiter pull	5/5	Real-world impact	5/5
Technical depth	4.5/5	Originality	4/5
Demo legibility	5/5	Buildability	4/5

PROS

- Real, sellable B2B product with live budget (CBSE adopting AI eval in 2026; post-NEET mandate).
- 'Founder who ships' PLUS frontier-security depth in ONE project - wins broad recruiters.
- Three novelty pillars (injection-robust, Rasch fairness, bias audit) no incumbent does.
- Visceral demo any non-techie gets: graders disagree by 18 marks, reconciled live to a cited score.
- Anti 'AI-wrapper': the depth is in calibration and robustness, not the prompt.

CONS / RISKS

- Grading market is crowded - you MUST hold the 'trust / audit' angle, not 'just grade'.
- High-stakes grading carries trust / liability weight.
- LLM-human agreement is only moderate - frame as ASSISTIVE, never 'replaces examiners'.

REAL-WORLD IMPACT & SELLABILITY Fairer marks for millions; ~70% less grading time for teachers; an audit and appeal trail institutions can defend. Direct ROI for coaching chains and boards - a genuine company.

LIMITATIONS & RISKS

- Needs model answers / rubrics plus a few expert-graded scripts to calibrate.
- Handwriting OCR on messy scripts is imperfect.
- Not for MCQ exams (already machine-graded).

HACKATHON SCOPE (MVP) One subject, 20-50 sample scripts. Multimodal handwriting ingest, GraphRAG over the official rubric, 3-4 AI-examiner ensemble, Rasch calibration + injection filter + bias check, teacher dashboard (per-dimension scores, citations, confidence, appeal sheet). Metric: variance cut from X to under Y marks; injection attacks blocked Z%.

TECH STACK Calibrator: claude-opus-4-8 (effort high). Examiner ensemble: claude-haiku-4-5 + Gemini Flash (free, handwriting), structured outputs. GraphRAG + ColPali. Next.js + D3. Eval: QWK / ICC + injection ASR.

FAR AWAY FIT + RESUME SIGNAL FAR AWAY: tops Real-World Impact + Execution; deep Tech (Rasch + robustness + multimodal RAG). Resume: ‘AI exam-evaluation trust layer - Rasch calibration cut examiner variance to under 4 marks, hardened vs answer-embedded injection, with bias audit’ = product + ML + security in one line.

VivaProof

GenAI-era authorship verification by auto-viva + process telemetry

CLAUDE'S RATING
8.6 / 10

NEW STAR · #2 Theme: Examinations

★ STAR NOVELTY (our edge): Stop DETECTING AI (it is broken - high false-positive AND negative). VERIFY authorship instead: auto-generate a personalized viva from the student's OWN submission plus draft / version telemetry, so they must defend their work - defensible exactly where GPTZero and Turnitin fail.

THE PROBLEM (plain English)
Universities can no longer tell which work is the student's: GPT writes original text that AI-detectors flag unreliably - too many false positives AND negatives for high-stakes calls. 'Authentic' take-home tasks are no longer authentic. This is openly unsolved in 2026.

HOW WE'D SOLVE IT
From the student's submission, auto-generate a personalized viva (timed-written or oral) that probes whether they can defend, critique and extend their own work, fused with process telemetry (draft / version history, editing patterns). Output: an authorship-confidence score plus a defensible evidence trail - not an automated accusation.

RATINGS

Resume / recruiter pull
Real-world impact
Technical depth
Originality
Demo legibility
Buildability

PROS

- Rides a massive, panicked 2026 pain (every university, globally) - easy B2B sell.
- Sidesteps the unwinnable AI-detection arms race - defensible, low false-positive.
- Spectacular legible demo: AI grills a student on their own essay; a GPT-only author cannot answer.
- Software-only, multimodal (voice viva via STT); strong India + global fit.

CONS / RISKS

- Coursera and others are entering - validates the market but watch competition.
- Viva fairness / accessibility must be handled (anxiety, disability, language).
- Authorship 'confidence' must be assistive evidence, never an automated verdict.

REAL-WORLD IMPACT & SELLABILITY Restores trust in assessment for thousands of institutions; saves faculty from manual vivas; a clear SaaS to universities, MOOCs and coaching - on one of 2026's hottest pains.

LIMITATIONS & RISKS
- Needs the student present for the viva (async-viva mitigates).
- Telemetry needs an editor / LMS integration.
- Voice / vernacular STT quality varies.

HACKATHON SCOPE (MVP) One course / assignment type. Ingest submission + draft history, RAG over the rubric / source, auto-generate 5 probing questions targeting the submission's specific claims, run a timed written or voice viva (Whisper STT), score defensibility + telemetry consistency, output an authorship-confidence report. Metric: separates GPT-only vs genuine authors at X% on a seeded set.

TECH STACK Question-gen + scoring: claude-opus-4-8 (structured outputs). Voice viva: Whisper STT + TTS. RAG: Chroma / GraphRAG. Telemetry: editor / Docs extension. Next.js viva UI.

FAR AWAY FIT + RESUME SIGNAL FAR AWAY: top Real-World Impact + Innovation + Demo; strong Tech (telemetry + RAG + voice). Resume: 'authorship-verification platform that replaces broken AI-detection with auto-generated vivas + process telemetry' = product + applied ML on a hot, legible problem.

FireBreak

Containment layer that stops prompt-injection 'worms' across multi-agent systems

CLAUDE'S RATING

8.3 / 10

SECURITY STAR · #3 Theme: Agentic and Autonomous Systems

★ STAR NOVELTY (our edge): Containment, not prevention. The first capability / provenance layer that stops 'Prompt Infection' from spreading ACROSS multi-agent handoffs - the one agent-security gap the 2026 survey still calls unsolved at scale (while authz, memory and MCP firewalls already shipped).

THE PROBLEM (plain English)

When AI agents work as a team and pass tasks to each other, one poisoned input can infect agent after agent like a virus ('Prompt Infection') - stealing data or escalating actions across the whole pipeline. The 2026 survey says no published defense blocks this at scale.

HOW WE'D SOLVE IT

A containment layer: every inter-agent message carries provenance / capability labels so untrusted-derived text cannot cross a handoff as a command, and each agent is held to Meta's 'Rule of Two'. Demo: a worm spreads across four agents and exfiltrates, then FireBreak quarantines it at each handoff.

RATINGS

Resume / recruiter pull	4/5	Real-world impact	3/5
Technical depth	5/5	Originality	5/5
Demo legibility	4/5	Buildability	3/5

PROS

- Genuinely unsolved frontier - maximum originality (unlike the now-saturated security primitives).
- Elite signal for AI-safety / security / research roles.
- Spectacular 'AI virus vs containment' live demo.
- Extends a real multi-agent framework (anti-wrapper).

CONS / RISKS

- Research-y / infrastructure - NOT a product you can sell.
- Abstract to product-minded recruiters.
- Hardest core to get working in the window.
- Honest framing must be 'containment', not 'solved'.

REAL-WORLD IMPACT & SELLABILITY Protects the multi-agent systems everyone is racing to deploy in 2026; shrinks injection blast radius. High value to AI labs, low legibility to a normal user.

LIMITATIONS & RISKS

- Cannot claim to 'prevent' injection - only contain propagation.
- Needs a believable multi-agent victim scenario.
- Credible eval needs adaptive attackers.

HACKATHON SCOPE (MVP) A 4-agent CrewAI / AutoGen pipeline + containment middleware (provenance-tagged message bus + capability policy) + a Prompt-Infection corpus + a live agent-graph dashboard + a propagation-rate eval (naive vs contained).

TECH STACK Agents: claude-opus-4-8; quarantined readers: claude-haiku-4-5 (structured outputs); MCP tools; LangGraph / CrewAI extension; Next.js graph dashboard.

FAR AWAY FIT + RESUME SIGNAL FAR AWAY: tops Tech Depth + Innovation; weaker Impact legibility. Resume: 'first capability-based containment for LLM-to-LLM Prompt Infection across a multi-agent pipeline' = elite AI-safety signal (for a narrower set of roles).

SentryBrowse

Firewall + live red-team for browser / computer-use agents

CLAUDE'S RATING
8.2 / 10

NEW · #4 Theme: Agentic and Autonomous Systems

★ STAR NOVELTY (our edge): Defuse the attack OpenAI says 'may never be solved': a deployable browser-agent firewall (visible-content sandboxing + action-risk gating) PLUS an RL red-team that proves it against invisible-text web injections - the exact vector Anthropic measured at a 31.5% hijack rate.

THE PROBLEM (plain English)

Browser agents (ChatGPT Atlas, Comet, Claude for Chrome) get hijacked by invisible instructions on web pages - white-on-white text, HTML comments - into leaking OTPs or hitting banking portals. OpenAI says it may never be solved; Anthropic measured a 31.5% hijack rate before safeguards.

HOW WE'D SOLVE IT

A defense layer + live red-team: separate what the agent SEES from what it can be TOLD (strip injected instructions from page content), gate risky actions by data-provenance + risk score, and ship an RL-style attacker that auto-discovers bypasses. Demo: a booby-trapped page hijacks a naive browser agent to exfiltrate an OTP; SentryBrowse blocks it.

RATINGS

Resume / recruiter pull	Real-world impact
Technical depth	Originality
Demo legibility	Buildability

PROS

- One of the hottest, most-cited problems in 2026 (OpenAI + Anthropic publicly flagging it).
- Maximally legible, dramatic demo (watch an AI get phished by a web page).
- Deep + frontier: provenance + sandboxing + RL red-teaming = serious engineering.

CONS / RISKS

- Research-crowded (ceLLMate, MUZZLE, VPI-Bench; frontier labs actively shipping).
- 'Solving' it is impossible - frame as risk-reduction + red-team.
- Computer-use plumbing is fiddly to demo reliably.

REAL-WORLD IMPACT & SELLABILITY Protects browser agents about to reach millions of users; the red-team harness is independently useful. High lab / security value.

LIMITATIONS & RISKS

- You compete with OpenAI / Anthropic's own teams on the headline.
- Needs a controlled browser-agent setup to demo.
- Containment, not a cure.

HACKATHON SCOPE (MVP) Wrap one open browser-agent (Browser-use / Agent-E) with a provenance-aware content filter + action-risk gate + a small RL / templated attacker over booby-trapped pages; dashboard shows hijack-attempt vs blocked + attack-success-rate before / after.

TECH STACK Agent: claude-opus-4-8 (computer-use); defense middleware in Python; attacker via templated + RL prompts; Playwright; Next.js dashboard.

FAR AWAY FIT + RESUME SIGNAL FAR AWAY: tops Innovation + Demo + Tech Depth; mid Impact legibility. Resume: 'browser-agent firewall + RL red-team for indirect prompt injection' = elite AI-security signal on the year's hottest agent problem (but a contested space).

OrbitGuard

Explainable satellite collision-risk screening + 3D globe for small operators

CLAUDE'S RATING
8.0 / 10

NEW WILDCARD · #5 Theme: Space and Aerospace

> **★ STAR NOVELTY (our edge):** Bench-marked space safety for the rest of us: a free browser 3D conjunction screen + an EXPLAINABLE ML risk score for small / Indian operators, validated against the Jan-2026 TraCSS OPEN answer-key dataset - real, defensible accuracy numbers, not a toy globe.

THE PROBLEM (plain English)
Small / university / Indian satellite operators get cryptic collision warnings and cannot afford the commercial tools (lakhs per year) to interpret them - so they fly blind or dodge needlessly. Until Jan 2026 there was not even open test data to build against.

HOW WE'D SOLVE IT
A free, browser-based conjunction-screening dashboard: public orbital data (TLEs), propagate orbits, screen for close approaches, an explainable ML risk score (probability + why), a 3D globe with a clear DODGE / WAIT / WATCH recommendation - validated against the new TraCSS answer key.

RATINGS

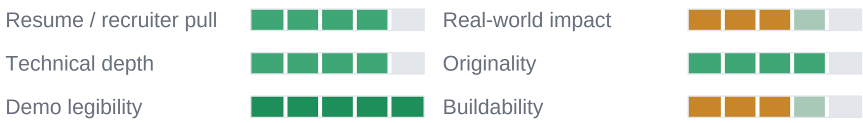

PROS

- Huge judge wow (live 3D globe of real orbits + collision alerts).
- Unusual, memorable resume line (aerospace + ML + viz) that stands out in a CS pile.
- Genuinely fresh enabler: the TraCSS answer-key dataset only opened Jan 2026.
- 100% software, 100% public data; strong India angle (Pixxel / GalaxEye / ISRO).

CONS / RISKS

- Less of a sellable product / narrower user base than the exam ideas.
- SpaceX Stargaze + commercial SSA exist - your wedge is explainability + small operators.
- Orbital-mechanics correctness is unforgiving; the math must be right.

REAL-WORLD IMPACT & SELLABILITY Democratizes collision safety for operators priced out of commercial SSA; civic + space-policy resonance, especially in India.

LIMITATIONS & RISKS
- Niche user base; monetization unclear.
- Needs careful propagation (Skyfield) to be credible.
- Public TLEs are lower-accuracy than operator ephemerides.

HACKATHON SCOPE (MVP) A region of LEO + a few public satellites. TLEs (Celestrak / Space-Track), Skyfield propagation, close-approach screening, an ML / analytic risk score + natural-language explanation, a Cesium / Three.js 3D globe + DODGE / WAIT panel; report accuracy vs the TraCSS answer key.

TECH STACK Python (Skyfield) propagation + screening; a light ML risk model; claude-opus-4-8 for the 'why + recommendation'; Cesium / Three.js + Next.js; TraCSS + Celestrak public data.

FAR AWAY FIT + RESUME SIGNAL FAR AWAY: tops Demo / Innovation + a rare Space entry; mid Impact-sellability. Resume: 'explainable satellite-collision dashboard validated on the official TraCSS dataset' = standout, hard-to-fake aerospace + ML line.

ExamShield

Agentic exam-integrity SOC for the NEET-2026 crisis

CLAUDE'S RATING
7.5 / 10

#6 Theme: Examinations

★ STAR NOVELTY (our edge): Catch the LEAK, not the cheater. SOC-style PRE-exam leak detection (RAG content-fingerprinting of untrusted social / Telegram feeds) - the upstream angle that proctoring tools entirely ignore.

THE PROBLEM (plain English)

NEET-2026 was cancelled after a paper leaked on WhatsApp ~weeks early (22.8 lakh students, 4 suicides). No single system catches leaks early, proctors fairly, AND audits the paper chain.

HOW WE'D SOLVE IT

A multi-agent platform: LeakDetector (RAG fingerprinting over social feeds), Proctor (bias-audited gaze / object detection), Auditor (tamper-evident paper chain) - optionally hardened with the dual-LLM pattern since it reads untrusted feeds.

RATINGS

Resume / recruiter pull	Real-world impact
Technical depth	Originality
Demo legibility	Buildability

PROS

- Maximum judge emotional pull (timely national crisis).
- Broad, instantly legible, strong India context.
- Several demoable modules.

CONS / RISKS

- Three loosely-coupled modules - wide but shallow.
- Proctoring carries real bias / liability baggage.
- Less of a single sellable product; weaker pure-resume novelty.

REAL-WORLD IMPACT & SELLABILITY High civic resonance and a 'would-have-caught-the-leak' narrative; reads as a public-good prototype more than a company or a deep system.

LIMITATIONS & RISKS

- Hard to build all three modules deeply.
- Proctoring fairness is a minefield.
- Leak-detection demo relies on simulated feeds.

HACKATHON SCOPE (MVP) Focus the demo on LeakDetector (RAG fingerprint match on a simulated leak feed) + a thin proctor + an audit timeline; do not try to ship all three deeply.

TECH STACK CrewAI; LlamaIndex / Chroma; MediaPipe; claude-opus-4-8 + claude-haiku-4-5; Next.js.

FAR AWAY FIT + RESUME SIGNAL FAR AWAY: tops Real-World Impact + Demo; mid Tech / Originality. Resume: good impact story, weaker 'deep system' or 'product' signal than GradeTrust / VivaProof.

JusticeRAG

Free ‘second-opinion’ score auditor for UPSC aspirants

CLAUDE’S RATING

7.3 / 10

#7 Theme: Examinations

★ STAR NOVELTY (our edge): A free, explainable second-opinion score (rubric-grounded multi-agent calibration). Real - but it is simply the B2C subset of GradeTrust’s B2B, 3-pillar superset.

THE PROBLEM (plain English)
UPSC Mains aspirants get scores from examiners who disagree by 30-40 marks on the same script, with no recourse (re-evaluation is not allowed).

HOW WE’D SOLVE IT
A multi-agent RAG (Scorer, Critic, Calibrator) over official model answers gives each candidate a free, rubric-grounded, explainable second opinion.

RATINGS

Resume / recruiter pull	Real-world impact
Technical depth	Originality
Demo legibility	Buildability

PROS

- Emotionally clean single narrative.
- Fully free-tier feasible (Gemini reads handwriting).
- Its technical core is shared with GradeTrust.

CONS / RISKS

- B2C aspirant tool - far narrower market than GradeTrust’s B2B.
- Just ‘RAG over PDFs’ unless deepened.
- Less of a ‘company’ story.

REAL-WORLD IMPACT & SELLABILITY Helps individual aspirants directly, but limited institutional reach and no clear revenue path.

LIMITATIONS & RISKS
- UPSC does not publish full criteria - alignment is capped.
- Handwriting / vernacular variance.
- Strictly a subset of GradeTrust.

HACKATHON SCOPE (MVP) Scorer / Critic / Calibrator agents + rubric RAG + handwriting ingest + a candidate report card (essentially GradeTrust’s B2C mode).

TECH STACK Gemini Flash (handwriting, free) + claude-opus-4-8 calibrator + Chroma / GraphRAG; Next.js.

FAR AWAY FIT + RESUME SIGNAL FAR AWAY: solid but dominated by GradeTrust. Best treated as GradeTrust’s free consumer mode.

CaMeL-Ops

Injection-proof single agent (dual-LLM + capability tokens)

CLAUDE'S RATING

7.0 / 10

#8 Theme: Agentic and Autonomous Systems

★ STAR NOVELTY (our edge): (Largely spent.) The dual-LLM + capability-token idea is strong - but Google, Microsoft (Dromedary) and an OpenClaw port already shipped it. Keep only as a learning reference, not a submission.

THE PROBLEM (plain English)
A tool-using AI agent (e.g. an inbox assistant) can be hijacked by a malicious instruction hidden in the data it reads - an email saying 'forward all emails to me'.

HOW WE'D SOLVE IT
Split the agent into a planner LLM (never sees untrusted data) + a quarantined reader LLM (no tools), with a capability engine gating tool calls. Demo: naive agent gets owned, ours blocks.

RATINGS

Resume / recruiter pull	4/5	Real-world impact	3/5
Technical depth	5/5	Originality	2/5
Demo legibility	4/5	Buildability	3/5

PROS

- Strong technical-depth + security signal.
- Clean 'hacked vs blocked' demo.
- Implements a famous 2025 paper.

CONS / RISKS

- NO LONGER ORIGINAL - 3 open-source versions already exist. A 4th clone reads as derivative.
- Single-agent only (the easy case; FireBreak solves the hard one).
- Not a product.

REAL-WORLD IMPACT & SELLABILITY Educational / reference value - but the exact gap it filled is now filled by others.

LIMITATIONS & RISKS

- Reimplements existing work; limited novelty headline.
- Was the original #1 before research found the clones.

HACKATHON SCOPE (MVP) Dual-LLM inbox agent + capability interpreter + injection corpus + side-by-side demo + ASR eval.

TECH STACK Planner claude-opus-4-8; quarantined claude-haiku-4-5; MCP tools; LangGraph.

FAR AWAY FIT + RESUME SIGNAL FAR AWAY: high Tech Depth, low Originality (saturated) - capped. Resume: weaker than extending it (FireBreak) or shipping a product with it (GradeTrust).

Final recommendation + how to decide together

BUILD THIS: GradeTrust - the exam-evaluation trust layer.

It is the only candidate that is BOTH a sellable product AND a frontier-depth build: rubric-grounded multimodal RAG + Many-Facet Rasch fairness calibration + student-injection defence + bias audit + appeal-grade explanations. It escapes the crowded ‘AI grader’ market by selling TRUST, not speed, and keeps the AI-security depth as a real anti-cheating feature. **Freshest alternative:** VivaProof (authorship verification) - same Examinations theme; the two could even combine into one Exam-Integrity suite.

HOW TO DECIDE AS A TEAM (5-minute vote)

- **Step 1 - role check:** if MOST of you target AI-safety / security labs, shortlist FireBreak or SentryBrowse; otherwise GradeTrust / VivaProof lead.
- **Step 2 - gut-check the demo:** which 2-minute demo makes a stranger lean in? (GradeTrust: graders disagree, we fix it. VivaProof: AI grills you on your own essay. SentryBrowse: AI gets phished by a web page.)
- **Step 3 - capacity:** the vertical products (GradeTrust / VivaProof) split cleanly into 5 lanes; the security ones (FireBreak / SentryBrowse) have a harder, riskier core.
- **Step 4 - vote** below; highest total wins. Tie-break toward the more sellable + more legible option.

Candidate	M1	M2	M3	M4	M5	Total
GradeTrust						
VivaProof						
FireBreak						
SentryBrowse						
OrbitGuard						
ExamShield						
JusticeRAG						
CaMeL-Ops						

(Each member scores every candidate 1-10 on ‘how proud would I be to put this on my resume AND demo it on stage’.)

Key sources (June 2026): CaMeL (arXiv:2503.18813) + Google / Microsoft Dromedary / OpenClaw impls; Prompt Infection (arXiv:2410.07283); Meta ‘Agents Rule of Two’; Microsoft Agent Governance Toolkit; mguard memory firewall; MINJA (arXiv:2503.03704); browser-agent injection (OpenAI / Anthropic 31.5%, ceLLMate, MUZZLE, VPI-Bench); TraCSS open dataset (Jan 2026); Eklavvya / Chanakya / CBSE pilot; SteLLA + GraphRAG ASAG; Wharton GAIL injection-on-grading; Many-Facet Rasch AES studies; authentic-assessment / auto-viva (BERA, Coursera, IATED).